

# UO Predictive Analytics

## 1 Bayesian Classification

This notebook is a demonstration of Naive Bayes classifier development for a loan approval prediction problem. We apply the concepts discussed: Bayes theorem and The Naive Bayes classifier

We will use the following python libraries for this practical. - Pandas: <https://pandas.pydata.org/pandas-docs/version/0.15/tutorials.html> - scikit-learn: <https://scikit-learn.org/stable/index.html>

Additionally, we will use the mixed-naive-bayes 0.0.1 library available at <https://pypi.org/project/mixed-naive-bayes/>.

**\*\*\* To install the package mixed-naive-bayes via pip in cmd.exe using the following:**

```
pip install git+https://github.com/remykarem/mixed-naive-bayes#egg=mixed_naive_bayes
```

If you are using Anaconda, please locate your Anaconda directory and use the command *activate* first.

**\*\*\*For Mac users, please install Anaconda3 and use the following command in Terminal:**

```
pip install mixed_naive_bayes
```

Similar to last week's practical, we will follow the workflow discussed in Concept 1.2 Predictive analytics workflow.

### 1.1 Predict Loan Eligibility for Dream Housing Finance company

source: <https://datahack.analyticsvidhya.com/contest/practice-problem-loan-prediction-iii/#ProblemStatement>

The dataset comes in csv format. Each sample has the following attributes:

Variable	Description
Loan_ID	Unique Loan ID
Gender_Married	Male/ Female
Dependents	Applicant married (Y/N)
Education	Number of dependents
Self_Employed	Applicant Education (Graduate/ Under Graduate) Self employed (Y/N)
ApplicantIncome	Applicant income
CoapplicantIncome	Coapplicant income
LoanAmount	Loan amount in thousands
Loan_Amount_Term	Term of loan in months
Credit_History	credit history meets guidelines
Property_Area	Urban/ Semi Urban/ Rural
Loan_Status	(Target) Loan approved (Y/N)

Our goal is to create a Naive Bayes model that can learn from the training samples, so that we can

predict outcome of a loan application.

```
[4]: #import the required libraries
import pandas as pd
from sklearn.model_selection import train_test_split
```

If you do not remember the purposes of the above imported modules. Please revisit week 1 practicals.

## 1.2 Step 1: Preprocessing

First, we load “loan\_prediction” dataset from the CSV file as pandas dataframe and observe first five instances.

```
[6]: df = pd.read_csv("loan_prediction.csv")
df.head(5)
```

```
[6]:
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	\
0	LP001003	Male	Yes	1	Graduate	No	
1	LP001005	Male	Yes	0	Graduate	Yes	
2	LP001006	Male	Yes	0	Not Graduate	No	
3	LP001008	Male	No	0	Graduate	No	
4	LP001011	Male	Yes	2	Graduate	Yes	

  

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	\
0	4583	1508.0	128.0	360.0	
1	3000	0.0	66.0	360.0	
2	2583	2358.0	120.0	360.0	
3	6000	0.0	141.0	360.0	
4	5417	4196.0	267.0	360.0	

  

	Credit_History	Property_Area	Loan_Status
0	1.0	Rural	N
1	1.0	Urban	Y
2	1.0	Urban	Y
3	1.0	Urban	Y
4	1.0	Urban	Y

```
[7]: #inspect the values in the categorical features
df['Gender'].value_counts()
```

```
[7]: Male      394
      Female    86
      Name: Gender, dtype: int64
```

```
[8]: df['Married'].value_counts()
```

```
[8]: Yes      311
      No       169
      Name: Married, dtype: int64
```

```
[9]: df['Dependents'].value_counts()
```

```
[9]: 0      274
      2      85
      1      80
      3+     41
      Name: Dependents, dtype: int64
```

Dependents is a categorical feature as 3+ is not a number. We should use it as a categorical feature.

```
[10]: df['Education'].value_counts()
```

```
[10]: Graduate      383
      Not Graduate   97
      Name: Education, dtype: int64
```

```
[11]: df['Self_Employed'].value_counts()
```

```
[11]: No      414
      Yes      66
      Name: Self_Employed, dtype: int64
```

```
[12]: df['Property_Area'].value_counts()
```

```
[12]: Semiurban    191
      Urban        150
      Rural        139
      Name: Property_Area, dtype: int64
```

Lets, take a quick look at the shape and summary of the dataset.

```
[13]: #shape
      df.shape
```

```
[13]: (480, 13)
```

```
[14]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 480 entries, 0 to 479
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Loan_ID               480 non-null   object
```

```

1  Gender          480 non-null  object
2  Married         480 non-null  object
3  Dependents      480 non-null  object
4  Education       480 non-null  object
5  Self_Employed   480 non-null  object
6  ApplicantIncome 480 non-null  int64
7  CoapplicantIncome 480 non-null float64
8  LoanAmount      480 non-null  float64
9  Loan_Amount_Term 480 non-null  float64
10 Credit_History  480 non-null  float64
11 Property_Area   480 non-null  object
12 Loan_Status     480 non-null  object
dtypes: float64(4), int64(1), object(8)
memory usage: 48.9+ KB

```

We get three important information from shape and summary: 1. There are 480 instances and 12 attributes. The target is Loan\_status. 1. From the Non-Null Count, we find that there is no missing values. Missing values affect the models adversely. We will learn about the effect of missing values and how to handle them later in the course. 1. dtypes at the bottom of the summary information tells us there are 4 floating point attributes, 1 integer attribute and 8 object or string valued attributes.

Similar to practical 1, we need to encode the categorical attributes.

```

[15]: #load the library for encoding
from sklearn import preprocessing

le = preprocessing.LabelEncoder()

df['en_gender']      = le.fit_transform(df['Gender'] )
df['en_married']     = le.fit_transform(df['Married'] )
df['en_dependents']  = le.fit_transform(df['Dependents'] )
df['en_education']   = le.fit_transform(df['Education'] )
df['en_self_employed'] = le.fit_transform(df['Self_Employed'] )
df['en_parea']       = le.fit_transform(df['Property_Area'] )

#encoding the target
df['target']         = le.fit_transform(df['Loan_Status'] )

```

```

[16]: #list the features
features = list(df.columns)
features

```

```

[16]: ['Loan_ID',
      'Gender',
      'Married',
      'Dependents',
      'Education',

```

```

'Self_Employed',
'ApplicantIncome',
'CoapplicantIncome',
'LoanAmount',
'Loan_Amount_Term',
'Credit_History',
'Property_Area',
'Loan_Status',
'en_gender',
'en_married',
'en_dependents',
'en_education',
'en_self_employed',
'en_parea',
'target']

```

We need to select the encoded features for our models. Also, loan id is an identifier i.e., it is different for each sample. We will exclude that from our model training.

```

[17]: #remove loan id and target - Loan_Status from features
features.remove('Loan_ID')
features.remove('Loan_Status')
features.remove('target')

#remove the non encoded features from the feature list
features.remove('Gender')
features.remove('Married')
features.remove('Dependents')
features.remove('Education')
features.remove('Self_Employed')
features.remove('Property_Area')

features

```

```

[17]: ['ApplicantIncome',
'CoapplicantIncome',
'LoanAmount',
'Loan_Amount_Term',
'Credit_History',
'en_gender',
'en_married',
'en_dependents',
'en_education',
'en_self_employed',
'en_parea']

```

```
[18]: #making sure we have 11 feautres in the list
len(features)
```

```
[18]: 11
```

Now, we check whether we have equal number of samples for each target class. A dataset where the samples are equally distributed to the target classes is called a balanced dataset. Imbalanced datasets are not good for training models as the model fails to perform well on samples who has the samples in the training set.

```
[19]: df['Loan_Status'].value_counts()
```

```
[19]: Y    332
      N    148
      Name: Loan_Status, dtype: int64
```

We observe that our data is not balanced. We have more samples where target class is Y than samples of class N.

Now we will divide our dataset into training and test sets. We will use a 70/30 split. Read about splitting using sklearn [https://scikit-learn.org/stable/modules/generated/sklearn.model\\_selection.train\\_test\\_split.html](https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.train_test_split.html)

As our dataset is imbalanced, we must use stratified sampling to split the dataset to ensure representative samples from all the target classes.

See the advantages of stratified sampling here [https://en.wikipedia.org/wiki/Stratified\\_sampling](https://en.wikipedia.org/wiki/Stratified_sampling).

```
[20]: #stratified sampling
X_train, X_test, y_train, y_test = train_test_split(df[features],          #the
                                                    df['target'],          #target
                                                    test_size=0.30,          #25% in
                                                    random_state=412,
                                                    stratify = df['target']
                                                    )
    ↪ feature space
    ↪ column
    ↪ the test set
    ↪ #the feautre to stratify
```

```
[21]: #check the class distribution in the test set
y_test.value_counts()
```

```
[21]: 1    100
      0    44
      Name: target, dtype: int64
```

### 1.3 Step 2: Learning the Naive Bayes model

To build a Naive Bayes classifier, we can use the implementation of the sklearn package. There are five different methods are provided for probability estimation. They are: 1. GaussianNB - for continuous or numeric attribute 1. MultinomialNB - for nominal or categorical attribute 1. ComplementNB - improved method for nominal or categorical attribute 1. BernoulliNB - for binary attributes 1. CategoricalNB - for nominal or categorical attribute

If our dataset has one type of features only i.e., datatypes of the features are only numeric or nominal then we can use one of the above implementations to build a predictive model. However, the features in the loan\_prediction dataset are of mixed types. Here, we cannot use sklearn.

The mixed-naive-bayes 0.0.1 (<https://pypi.org/project/mixed-naive-bayes/#quick-start>) provides an implementation of Naive Bayes for mixed attributes. We will use this python library for this practical.

Note: The module expects that we have label encoded the categorical features.

```
[22]: #initialize the mixed_naive_bayes lib
      from mixed_naive_bayes import MixedNB

      #fit the model with the training set
      clfNB = MixedNB(categorical_features=[5,6,7,8,9,10])
      #categorical_features is a list of indices categorical attributes in our dataset

      clfNB.fit(X_train, y_train)
```

```
[22]: MixedNB(alpha=0.5, var_smoothing=1e-09)
```

### 1.4 Step 3: Evaluation

We have trained our model in the previous step which is represented as knn. Read about sklearn kNN method here <https://scikit-learn.org/stable/modules/generated/sklearn.neighbors.KNeighborsClassifier.html>.

We will now evaluate the performance of our model on the test set. That is, we will apply the model to the test set X\_test and match the predictions of the model with y\_test.

```
[23]: #predictions
      y_pred = clfNB.predict(X_test)
      y_pred
```

```
[23]: array([1, 1, 0, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
         0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 1, 1, 1, 1, 0, 0, 1, 1, 1, 0, 1, 1,
         0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 0, 1, 1, 1, 1, 1,
         1, 1, 1, 1, 1, 1, 0, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 0, 1, 1, 1, 1,
         0, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 0, 1, 1, 1, 0, 1, 0, 0, 1, 0, 1,
         1, 0, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 0, 1, 1, 0, 1, 0, 1, 0, 1, 1,
         0, 0, 0, 0, 1, 1, 1, 1, 1, 1, 0, 1], dtype=int64)
```

```
[24]: #here is our true labels
      y_test
```

```
[24]: 421    1
      394    1
      347    1
      317    1
      454    0
      ..
      17    1
      31    1
      72    1
      445    0
      307    1
      Name: target, Length: 144, dtype: int32
```

Detailed list of parameters of MixedNB is provided here [https://remykarem.github.io/docs/mixed\\_naive\\_bayes.html](https://remykarem.github.io/docs/mixed_naive_bayes.html)

For the test instances, we can inspect the class probabilities as well.

```
[25]: #let us take the first test instance for example.
      X_test[0:1]
```

```
[25]: ApplicantIncome  CoapplicantIncome  LoanAmount  Loan_Amount_Term \
421              6417              0.0          157.0          180.0

      Credit_History  en_gender  en_married  en_dependents  en_education \
421              1.0          1          1          3          0

      en_self_employed  en_parea
421              0          0
```

```
[26]: clfNB.predict_proba(X_test[0:1])
```

```
[26]: array([[1.91465654e-11,  2.81453864e-11]])
```

#### 1.4.1 Explanation

Given an instance  $\mathbf{x}' = (6417, 0.0, 157.0, 180.0, 1.0, 1, 1, 3, 0, 0, 0)$ , the model predicts

$$P(Y|\mathbf{x}') = 1.915 \times 10^{-11}$$

$$P(N|\mathbf{x}') = 2.815 \times 10^{-11}$$

We can say that the model predicted that the application is not eligible for a loan.

```
[27]: #measure accuracy
      from sklearn.metrics import accuracy_score
```



```
accuracy_score(y_test, y_pred)
```

[27]: 0.8125

```
[28]: #we can also use the score method of MixedNB to observe the accuracy  
clfNB.score(X_test, y_test)
```

[28]: 0.8125

Our Naive Bayes model has 81.25% accuracy on the loan\_prediction dataset.

In this practical, we work through all the steps required to develop a Naive Bayes model. We used a 70/30 splits for training and test sets.